

OPENAI / DFIR / CHEAT SHEET



invictus-ir.com | info@invictus-ir.com

A compact guide to what could exist, what to collect, and what the evidence can establish across the ChatGPT Enterprise workspace, the API Platform organisation, and agentic workloads (Codex, Agents SDK).

01 SCOPE AND CONTROLS FIRST

ENVIRONMENT
ChatGPT workspace and API Platform organisation are separate evidence environments. Record the workspace_id, organization_id, every project_id, then users and service accounts.

LOGGING STATE
Confirm whether API Platform audit logging was on, and since when. Record api_call_logging and its project scope. Neither is retroactive.

DATA CONTROLS
Capture ZDR and Modified Abuse Monitoring status, workspace retention and residency. Under ZDR, stored response content and hosted Agents SDK tracing are unavailable.

TOOL AND NETWORK ACCESS
Record whether MCP, web search, file search, code interpreter and container outbound network were disabled or limited to selected projects.

02 EVIDENCE MAP

CHATGPT ENTERPRISE / EDU WORKSPACE · via the Compliance Platform			API PLATFORM ORGANISATION · via the Admin API		
SOURCE	WHAT IT GIVES YOU	RETENTION	SOURCE	WHAT IT GIVES YOU	RETENTION
Compliance Logs Platform	Prompt and completion content, uploaded files, GPT config, memories, users. Content follows the workspace retention policy.	30 days	API Platform Audit Logs /v1/organization/audit_logs	Identity, credential, role, invite, project, login, org config change. Enable in advance, cannot disable once on.	Best effort
Stateful Compliance API	Current object state and conversations retained by the workspace. Route beyond the 30-day log window.	Per policy	API Call Logs	Request and response content for supported calls. api_call_logging must cover the project.	Per Call
Admin Audit Logs	Workspace and administrative change. Diff against a known-good baseline.	30 days	Stored Responses and Conversations	Retained model input and output, conversation state. ZDR forces store off for eligible endpoints.	30 days
User Authentication Logs	Login and access activity, source context. Correlate with the IdP.	30 days	Agents SDK Traces	Generations, tool and function calls, handoffs, guardrails, span tree. Client-side and easily suppressed.	Not fixed
Codex Usage Logs	Coding-agent activity on supported surfaces. Coverage is surface-specific.	30 days	Files and Vector Stores	Object existence, content, upload and attachment lineage. Often outlive logs.	Until deleted

03 INVESTIGATION PATHS

SCENARIO	COLLECT	EVENTS / EVIDENCE	ESTABLISHES
Account compromise Did an actor gain access, and what did they do next?	User Authentication Logs Compliance Logs Platform Activity after authentication	login.succeeded login.failed logout.* user.* project.* api_key.* role activity	- Authentication timeline - Actor and session context - Administrative activity that followed
Persistence and privilege escalation Did the actor build durable access or widen privileges?	Admin Audit Logs Workspace access and administrative change	api_key.* service_account.created (role=owner) invite.* user.* role.updated role.assignment.created	- New credentials - Service accounts - Membership changes - Roles and permission grants
Defence evasion Did the actor weaken logging, identity or network controls?	Admin Audit Logs Security and administrative change	organization.updated (api_call_logging narrowed) ip_allowlist.config.deactivated scim.disabled certificate.* external_key.removed	A control moving to off or reduced coverage
Data exposure What data reached OpenAI, and what did the model receive or return?	Compliance Logs Platform Retained files GPTs Memories	API Call Logs Stored Responses and Conversations Files and Vector Stores Audit Logs for actor and key	- Recorded input and output - Retained workload state
Agent and Codex abuse What did the agent attempt, who initiated it, what executed?	Codex Usage Logs Compliance Logs Platform Retained artefacts	Agents SDK Traces API Call Logs Audit Logs for actor and key	- Generations - Tool and function calls - Handoffs